

FEWS: Large-Scale, Low-Shot Word Sense Disambiguation with the Dictionary

Terra Blevins, Mandar Joshi, and Luke Zettlemoyer

Paul G. Allen School of Computer Science & Engineering, University of Washington

Facebook AI Research, Seattle

{blvns, mandar90, lsz}@cs.washington.edu

Abstract

noncurrent models for Word Sense Disambiguation (WSD) struggle to disambiguate rare senses, despite reaching nonhuman performance on global WSD metrics. This stems from a lack of data for both modeling and evaluating rare senses in existing WSD datasets. In this paper, we introduce FEWS (many-shot Examples of Word Senses), a new high-shot WSD dataset automatically extracted from example sentences in Wiktionary. FEWS has

low spirit sense coverage across like supernatural language domains and provides:
(1) a small training set that covers few less senses than previous datasets and (2) a noncomprehensive evaluation set containing numbers or terms having various uses and zero-shot examples of a narrow variety of senses. We establish baseline on FEWS with knowledge-based and neural WSD approaches and future transfer learning experiments demonstrating that models additionally trained on shot WSD.

1 Introduction

Word Sense Disambiguation (WSD) is the task of identifying the *sense*, or meaning, that an ambiguous word takes in a specific context. Recent WSD models (Huang et al., 2019; Blevins and Zettlemoyer, 2020; Bevilacqua and Navigli, 2020) have made large gains on the task, surpassing the estimated 80% F1 human performance upper bound on WordNet annotated corpora (Navigli, 2009). Despite this breakthrough, the task remains far from solved: performance on rare and zero-shot senses is still low, and in general, current WSD models struggle to learn senses with few training examples (Kumar et al., 2019; Blevins and Zettlemoyer, 2020). This performance gap stems from limited data for rare senses in current WSD datasets, which are annotated on natural language documents that contain a Zipfian distribution of senses (Postma et al., 2016).

More generally, since each word has a different set of candidate senses and new senses are regularly coined, it is almost impossible to gather a large number of examples for each sense in a language.

C: I **liked** my friend’s last status...

S1: to enjoy... [or] be in favor of.

✓ S2: To show support for, or approval of, something on the Internet by marking it with a vote.

C: A transistor-diode **matrix** is composed of vertical and horizontal wires with a transistor at each intersection.

✓ S1: A grid-like arrangement of electronic components, especially one intended for information coding, decoding or storage.

S2: A rectangular arrangement of numbers or terms having various uses in mathematics.

Figure 1: Sample contexts (C) from FEWS with ambiguous words and a subset of candidate sense definitions (S). FEWS covers a wide range of senses, including new senses and domain-specific senses.

This makes the few-shot learning setting particularly important for WSD. We introduce **FEWS** (Few-shot Examples of Word Senses), a dataset built to comprehensively train and evaluate WSD models in few- and zero-shot settings. Overall, the contributions of FEWS are two-fold: as training data, it exposes models a broad array of senses in a low-shot setting, and the large evaluation set allows for more robust evaluation of rare senses.

FEWS achieves high coverage of rare senses by automatically extracting example sentences from Wiktionary definitions. Wiktionary is an apt data source for this purpose, containing examples for over 71,000 senses (Table 1). Not only is this sense coverage higher than existing datasets (e.g., SemCor, the largest manually annotated WSD dataset, only covers approximately 33,000 senses (Miller et al., 1993)), it also extends to senses related to new domains (Figure 1).

Data Split	# Examples	# Tokens	# Annot.	# Sense Types	# Word Types	Ambiguity
Overall	121,459	3,259,240	131,278	71,391	35,416	5.62
Train	87,329	2,551,358	96,023	52,928	30,450	4.98
Ext. Train	101,459	2,683,345	111,278	61,391	31,937	5.71
Dev	10,000	287,673	10,000	10,000	8,682	5.09
Few-shot	5,000	149,791	5,000	5,000	4,417	4.77
Zero-shot	5,000	137,882	5,000	5,000	4,661	5.41
Test	10,000	288,222	10,000	10,000	8,709	5.10
Few-shot	5,000	149,384	5,000	5,000	4,449	4.71
Zero-shot	5,000	138,838	5,000	5,000	4,666	5.49

Table 1: FEWS data statistics. The development and test sets are balanced across senses and split evenly between few-shot examples (with support in the training set) and zero-shot examples. The extended training set (Ext. Train) adds short examples written by Wiktionary editors as additional training data.

We establish performance baselines on FEWS with both knowledge-based approaches and a recent neural biencoder model for WSD (Blevins and Zettlemoyer, 2020). We find that the biencoder, despite being the strongest baseline on FEWS, still underperforms human annotators, particularly on zero-shot senses where the biencoder trails by more than 10%. We also present transfer learning experiments and find adding FEWS as additional training data improves performance on all but the most frequent senses (MFS) in the WSD Evaluation Framework (Raganato et al., 2017); this suggests that future improvements on FEWS could generalize other WSD benchmarks. FEWS is available at <https://nlp.cs.washington.edu/fews>.

2 Related Work

WSD is a long-standing NLP task and is the focus of many datasets. The current *de facto* benchmark for modeling English WSD is the WSD Evaluation Framework (Raganato et al., 2017), which includes the SemCor dataset (Miller et al., 1993) as training data and consolidates a number of evaluation sets (Pradhan et al., 2007a; Palmer et al., 2001; Snyder and Palmer, 2004; Navigli et al., 2013; Moro and Navigli, 2015) into a standardized evaluation suite. These datasets are annotated with the senses (known as *synsets*) from Wordnet, a manually constructed ontology of semantic relations (Miller et al., 1993).

Most existing datasets for WSD, including those in the WSD Evaluation framework and others like Pradhan et al. (2007b), are annotated on natural language documents that contain a Zipfian distribution of word senses (Kilgariff, 2004). This data source causes these datasets to have low coverage of rare senses, leading to worse performance on these less common senses (Postma et al., 2016; Kumar et al.,

2019). In contrast, we use example sentences from Wiktionary as an alternative source of text for WSD data with FEWS. This means that FEWS has a more uniform sense distribution, providing more balanced coverage across different senses of words.

Wiktionary has previously been used as a resource for WSD research. Most work has investigated mapping Wiktionary senses onto WordNet synsets (Meyer and Gurevych, 2011; Matuschek and Gurevych, 2013); other work has learned similar mappings for Wikipedia articles (Mihalcea (2007); Navigli and Ponzetto (2012); inter alia). More similar to our work, Henrich et al. (2012) and Segonne et al. (2019) mine WSD examples from Wiktionary to augment labeled WSD data for non-English languages. However, FEWS is the first dataset specifically designed to evaluate zero and few-shot learning with the balanced dictionary sense distribution.

3 FEWS: Low-shot Learning for WSD

FEWS (Few-shot Examples of Word Senses) is a new dataset for learning to do low-shot WSD.¹ It is created with example contexts drawn from Wiktionary, an online collaborative dictionary.² Since Wiktionary is curated by volunteers, the data is manually annotated and high quality, and there is no additional annotation cost to construct FEWS. Furthermore, using example contexts from a dictionary allows FEWS to cover many senses under-represented in existing WSD datasets, such as rare senses of words or senses pertaining to specific domains. However, we note that due to the crowd-sourced nature of the examples in Wiktionary and

¹We use the term *low-shot* as an umbrella term for few- and zero-shot learning.

²<https://en.wiktionary.org/>




Figure 2: Comparison of sense coverage for five words in the Semcor and FEWS training corporuses.

the subjectivity of fine-grained sense distinctions, inconsistencies in the underlying data may introduce some annotation errors into FEWS.

Examples of the data in FEWS are shown in Figure 1. In FEWS, each example context contains one or more instances of the ambiguous *target word* and is labeled with the sense (and corresponding definition) that describes that word as used in the context; this is in contrast to *all-words WSD*, where many of the content words in the context are annotated.

3.1 Dataset Creation

To create FEWS, we extracted all of the definitions for content words (nouns, verbs, adjectives, and adverbs) and example contexts associated with those definitions from a checkpointed version of English Wiktionary.³ While processing the Wiktionary data, we collected two types of contexts: (1) *quotations* (93% of extracted examples), which are quotations of natural language text found by Wiktionary contributors that contain the target word used with the relevant sense, and (2) *illustrations* (7%), which are short sentences or fragments written by contributors to illustrate the word sense in context. The target words in each context are marked by the Wiktionary formatting metadata; examples where no words are marked or the marked word differs too much from the base form of the dictionary entry are discarded.⁴ We additionally filter out examples that are too short to provide a meaningful context for the marked word.

³We focus English senses and filter out definitions for words in other languages; however, this data collection process could be expanded to other languages in Wiktionary.

⁴We define a marked word as too different from the base form if the longest common subsequence between them is < 50% of the length of the marked word.

We then labeled the target words in each extracted context with the sense ID generated for the definition associated with that sentence; this gave us 254,506 annotated WSD example contexts covering 148,333 senses. Finally, we filtered out examples with monosemous words since predicting the sense in these cases is a trivial task. However, Loureiro and Camacho-Collados (2020) recently found that unambiguous examples can improve WSD performance; therefore, we include the unambiguous cases as an additional file in the FEWS dataset. After filtering these unambiguous examples from the main dataset, FEWS in total contains 121,459 examples covering 71,391 sense types.

Finally, we split the data into training and evaluation sets. The majority of the data are *quotations*, which we use to populate the train, development, and test sets as they more closely resemble naturally occurring text than the *illustrations*. To create the development and test sets, we randomly select 10,000 examples for each evaluation set and ensure that each of these evaluation examples pertains to a different sense. We verify that half of those examples were labeled with senses that only occurred once in the unsplit data to create a zero-shot subset of each evaluation set, and the other half of the evaluation senses comprise the few-shot evaluation subsets. The remaining quotations that were not used for the development or test set are included as the training data. Finally, we remove the *illustrations* for senses in the zero-shot evaluation subsets and add the remaining *illustrations* to the training data; this addition makes the extended train set.

3.2 Dataset Analysis

We present a comprehensive analysis of FEWS to demonstrate that the dataset provides high coverage

internet culture.

We also find that a number of the senses (<1% in FEWS) are indicated to be toxic or offensive language: our analysis contains tags such as *ethnic slurs*, *offensive*, *vulgar*, and *derogatory* that correspond to examples of toxic language. For many of these examples, the meaning of the target word is only toxic due to the context in which it appears. These examples provide an opportunity for improving models for hate speech detection and related tasks, but we leave this exploration to future work.

4 Baselines for FEWS

We run a series of baseline approaches on FEWS to demonstrate how current methods for WSD perform on this dataset. We consider a number of knowledge-based approaches (**Lesk**, **Lesk+emb**, and **MFS**) and two neural models that build on pre-trained encoders (**Probe_{BERT}** and **BEM_{BERT}**). We also ascertain how well humans perform on FEWS as a potential upper bound for model performance.

4.1 Knowledge-based Baselines

Most Frequent Sense (MFS) The MFS baseline assigns each target word in the evaluation with their candidate sense that is most frequently observed as the correct sense of that word in the training set. The MFS heuristic is known to be a particularly strong baseline in WSD datasets labeled on natural language documents (Kilgarriff, 2004); however, we expect this to be a weaker baseline on FEWS since the distribution of senses is much more balanced (and half of the evaluation senses are completely unseen during training).

Lesk The simplified **Lesk** algorithm assigns to each ambiguous target word the sense whose gloss has the highest word overlap with the context surrounding that target word (Kilgarriff and Rosenzweig, 2000). We specifically use the Lesk-definitions baseline from this work, meaning that we do *not* include words from example sentences in the set compared against the context – since these example sentences are used as the contexts in FEWS.

Lesk+emb This baseline is an extension of the above approach that incorporates word embeddings (Basile et al., 2014). A vector representation is built for the context around an ambiguous word (v_c) and the glosses of each sense of that word

(v_g), where v_c and v_g are the element-wise sums of the word vectors for words in the context and gloss, respectively. The sense that corresponds to the v_g with the highest cosine similarity to v_c is chosen as the label for the target word. We use Glove embeddings (Pennington et al., 2014) for our implementation of this baseline.

4.2 Neural Baselines

Probe_{BERT} This baseline is a linear classifier trained on contextualized representations output by the final layer of a frozen pretrained model; we use BERT as our pretrained encoder (Devlin et al., 2019). We train this classifier by performing a softmax over all of the senses in the Wiktionary sense inventory and mask out any senses not relevant to the target word.

BEM Our other neural baseline is the biencoder model (BEM) for WSD introduced by Blevins and Zettlemoyer (2020). The BEM has two independent encoders, a *context encoder* that processes the context sentence (including the target word) and a *gloss encoder* that encodes the glosses of senses into a sense representation. The BEM takes the dot product of the target word representation from the *context encoder* and sense representations from the *gloss encoder*, and it labels the target word with the sense that has the highest dot product score. We train **BEM_{BERT}** by initializing each encoder with BERT and training on the FEWS train set.

4.3 Human Performance

Finally, we calculate the estimated human performance on the FEWS development set. The three human annotators were native English speakers, who each completed the same randomly chosen 300 example subset of the development set. The examples were sampled such that half (150) of these examples came from the few-shot split and the other half came from the zero-shot split. Similar to the modeling baselines, we evaluate each annotator’s performance by scoring them against the sense associated with that example in Wiktionary (which we assume to be gold labels).

5 Baseline Experiments

5.1 Experimental Setup

Data All baselines for the FEWS dataset are trained on the *train* set unless specifically stated to have been trained on the *extended train* set. All

	Dev			Test		
	Full Set	Few-shot	Zero-shot	Full Set	Few-shot	Zero-shot
Knowledge-based baselines						
MFS	26.39	52.78	0.00	25.73	51.46	0.00
Lesk	39.24	42.54	35.94	39.07	40.94	37.20
Lesk+emb	42.54	44.94	40.14	41.53	44.08	38.98
Human [†]	80.11	80.44	79.87	–	–	–
Neural baselines						
Probe _{BERT}	36.17	72.34	0.00	36.07	72.14	0.00
BEM _{BERT}	73.81	79.28	68.34	72.77	79.06	66.48
BEM _{SemCor}	74.36	79.72	69.00	72.98	78.88	67.08
BEM _{zero-shot} [‡]	58.05	58.38	57.72	57.39	57.94	56.84

Table 3: Accuracy (%) of our baselines (Section 4) and transfer learning models (Section 6) on the FEWS evaluation sets. Human performance ([†]) is calculated on a subset of the development set and acts as an estimated upper bound on performance. **Probe_{BERT}** and **BEM_{BERT}** are baselines trained on FEWS; **BEM_{SemCor}** is a transfer learning model finetuned on SemCor before training on FEWS. **BEM_{zero-shot}** ([‡]) is a zero-shot transfer experiment in which the BEM trained on SemCor is evaluated on FEWS *without* finetuning on the FEWS train set.

models for FEWS are tuned using the *development* set and then evaluated on the held-out *test* set.

Experimental Details Our probe and BEM baselines are implemented in PyTorch⁵ and optimized with Adam (Kingma and Ba, 2015). For the BEM, we use the implementation provided by Blevins and Zettlemoyer (2020).⁶ We obtain the *bert-base-uncased* encoder from Wolf et al. (2019) to get the BERT output representations for the probe and to initialize the BEM models. Further hyperparameter details are given in Appendix A.

5.2 Modeling Results

Table 3 shows the results of our baseline experiments on FEWS. We find that the MFS baseline is weak overall, primarily because it is unable to predict anything about the held-out, zero-shot senses; the Lesk algorithms both outperform this baseline in the overall setting, with the Lesk+emb approach scoring slightly better than the original Lesk approach by 1.78-4.2% across the different evaluation subsets. However, on the few-shot examples in both the development and test sets, we see that the MFS baseline outperforms both of the Lesk baselines. This shows that, for the few-shot examples, the MFS heuristic remains a reasonably strong baseline even with the more uniform sense distribution of FEWS (and indicates that the distribution of examples drawn from the dictionary is

less uniform than expected).

The neural baselines we run generally outperform the knowledge-based ones. The **Probe_{BERT}** model does fairly well on the few-shot examples, outperforming the MFS baseline by about 20 accuracy points; however, it is unable to disambiguate words in the zero-shot splits correctly since the probe can not generalize to unseen senses. In comparison, **BEM_{BERT}** performs well across the entire evaluation set. In particular, the BEM achieves much better zero-shot performance than other baselines, though performance on this subset still lags behind few-shot performance. Finally, we see that humans perform better than all of the considered baselines, particularly on zero-shot senses where humans outperform the BEM by 11.53 points. More details about the human evaluation are given below (Section 5.3).

Additionally, we find that training on the extended train set has little effect on these baselines: the **MFS** and **Probe_{BERT}** baselines perform slightly worse (with deltas of -0.08% and -0.21% on the test set, respectively), and **BEM_{BERT}** performs 1.05% better. Appendix B presents full results for the extended train baselines.

5.3 Human Evaluation Results

The human annotators achieved an average accuracy of 80.11% (with each annotator getting 84.67%, 78.67%, and 77.00% accuracy) and an average inter-annotator agreement of $\kappa = 0.802$. We find that humans perform slightly better on the

⁵<https://pytorch.org/>

⁶<https://github.com/facebookresearch/wsd-biencoders>

	Dev Subset		
	Full Set	Few-shot	Zero-shot
Knowledge-based baselines			
MFS	23.66	47.33	0.00
Lesk	38.33	41.33	35.33
Lesk+emb	46.00	48.67	43.33
Human	80.11	80.44	79.87
Neural baselines			
Probe _{BERT}	33.67	67.33	0.00
BEM _{BERT}	73.00	80.66	65.33

Table 4: Accuracy (%) of our baselines on the subset of the development set manually scored by human annotators.

examples that correspond to few-shot examples for the dataset than those corresponding to zero-shot examples despite not using the training data, with an average of 80.44% and 79.87% accuracy on those two subsets, respectively.

We also report the performance of the baselines for FEWS on the same subset of the development set that was manually completed by humans (Table 4). We find that the baselines perform similarly on this subset, with a small decrease in performance compared to the full development set (with decreases in accuracy ranging between 0.81% and 3.46% when compared to the development set).

6 Transfer Learning with FEWS

Next, we investigate how useful FEWS is at improving WSD performance on existing benchmarks. We perform transfer learning experiments by iteratively finetuning models on FEWS and the WSD Evaluation Framework (Raganato et al., 2017), with one acting as the *intermediate* dataset and evaluating performance on the other, *target* dataset. We find that on global metrics, this approach performs similarly to finetuning only on the training data for each benchmark; however, transferring from FEWS to the WSD Framework improves performance on less-frequent and zero-shot senses. This suggests that FEWS provides valuable WSD information not covered in SemCor.

6.1 Experimental Setup

We apply the supplementary training approach presented in Phang et al. (2018) to perform our transfer learning experiments. We initialize a BEM with the best model developed on the *intermediate* dataset

and evaluate on the *target* dataset in two ways: first, by evaluating this BEM on the *target* dataset in a zero-shot manner, without additional finetuning; and second by finetuning the BEM on the *target* training data before performing the *target* evaluation. We refer to these models as **BEM_{zero-shot}** and **BEM_{intermediate}**, respectively. As a baseline, we also compare against the best BEM trained only on the *target* dataset (with no exposure to the *intermediate* dataset), **BEM_{BERT}**.⁷ The model implementation details are identical to those for the baseline experiments on FEWS (Section 5.1).

6.2 Data

Models that are finetuned for the WSD Evaluation Framework are trained using SemCor, a large dataset annotated with WordNet synsets and commonly used for training WSD models (Miller et al., 1993). Following previous work, we use SemEval-2007 as a validation set (SE07; (Pradhan et al., 2007a)) and hold out the other evaluations sets in the Framework (Senseval2 (SE2; (Palmer et al., 2001)), Senseval-3 (SE3; (Snyder and Palmer, 2004)), SemEval-2013 (SE13; (Navigli et al., 2013)), and SemEval-2015 (SE15; (Moro and Navigli, 2015))) as test sets. Similarly to the baseline experiments, models that are trained on FEWS use the *train* set (note that we do not use the *extended train* set in these experiments), are validated on the *development* set, and finally evaluated on the held out *test* set.

6.3 Results

FEWS Results We first consider the setting where SemCor acts as the *intermediate* dataset and FEWS as the *target* (Table 3). We find that **BEM_{SemCor}** performs similarly to training only on FEWS (with an improvement of 0.21% on the test set). **BEM_{zero-shot}**, which is not finetuned on the FEWS train set, unsurprisingly performs worse than any of the BEMs that saw the FEWS training data but outperforms the Lesk baselines.

WSD Evaluation Framework Results We then consider the opposite setting, where FEWS is the *intermediate* dataset and the WSD Evaluation Framework acts as the *target* evaluation (Table 5).

⁷We note that this is a naive approach to transferring between Wiktionary senses and WordNet synsets, and that it is possible that better transfer learning could occur with a more complicated multitask approach or by mapping between the two lexical resources with approaches such as Navigli and Ponzetto (2012) or Miller and Gurevych (2014).

	Dev	Test Datasets					Concatenation of all Datasets				
	SE07	SE2	SE3	SE13	SE15		Nouns	Verbs	Adj.	Adv.	ALL
MFS (in train data)	54.5	65.6	66.0	63.8	67.1		67.7	49.8	73.1	80.5	65.5
BEM _{BERT}	74.5	79.4	77.4	79.7	81.7		81.4	68.5	83.0	87.9	79.0
BEM _{FEWS}	73.6	79.1	77.9	79.1	81.6		81.2	68.9	81.8	88.2	78.8
BEM _{zero-shot}	53.0	66.6	62.3	69.1	74.9		70.7	51.2	72.0	69.7	66.4

Table 5: F1-score on the English all-words WSD in the WSD Evaluation Framework (Raganato et al., 2017). We compare the best model from Blevins and Zettlemoyer (2020) (BEM_{BERT}) against (1) a model first trained on FEWS and then trained on SemCor (BEM_{FEWS}) and (2) a model trained on FEWS and evaluated on this task without further finetuning (BEM_{zero-shot}).

	MFS	LFS	Zero-shot	
			Words	Senses
WordNet S1	100.0	0.0	84.9	53.9
BEM _{BERT}	94.1	52.6	91.2	68.9
BEM _{FEWS}	93.7	52.9	92.2	74.8
BEM _{zero-shot}	72.6	55.5	92.7	80.5

Table 6: F1-score on the MFS, LFS, and zero-shot subsets of the **ALL** evaluation set from the WSD Evaluation Framework. Zero-shot examples are the words and senses (respectively) from the evaluation suite that do not occur in SemCor.

On the overall evaluation set, we again see that BEM_{FEWS} performs similarly to the BEM_{BERT} baseline on the overall **ALL** metric, and that the zero-shot BEM model underperforms the other bi-encoders.

We then break down performance on the *target* evaluation set by sense frequencies: we evaluate performance on the most frequent sense (MFS) of each word in the evaluation (i.e., the sense each word takes most frequently in the SemCor training set); the less frequent senses (LFS) of words, or any sense a words takes besides its MFS; zero-shot words that are not seen during training on SemCor; and zero-shot senses, also not seen during training (Table 6). We find that both transfer learning models perform better on LFS and zero-shot examples than BEM_{BERT}.⁸

In particular, the zero-shot transfer model does well on these subsets, demonstrating that a fair amount of WSD knowledge about uncommon senses can be transferred in a zero-shot manner between these datasets (albeit at the expense of higher

performance on the MFS subset). This result also shows how much the MFS group dominates existing WSD metrics and highlights the need for focused evaluations of other types of word senses. Finally, we see that even without exposure to the natural sense distribution in natural language texts, the zero-shot model still performs significantly better on the MFS of words than the LFS, with a 17.1 F1 point difference between the two subsets; this is likely because the BERT encoder is exposed to the sense distribution of English natural language documents during pretraining.

7 Conclusion

We establish baseline performance on FEWS with both knowledge-based approaches and recently published neural WSD models. Unsurprisingly, neural models based on pretrained encoders perform best on FEWS; however, the human evaluation shows there is still room for improvement, particularly for zero-shot senses. Finally, we also present results on transferring word sense knowledge from FEWS onto existing WSD datasets with staged finetuning. While our naive approach for transferring knowledge from FEWS does not improve performance on the global WSD metric, adding FEWS as an additional training signal improves performance on uncommon senses in existing evaluation sets.

We hope that FEWS will inspire future work focusing on better methods for capturing rare senses in WSD and better modeling of word sense in niche domains like internet culture or technical writing.

Acknowledgements

This material is based on work conducted at the University of Washington, which was supported by the National Science Foundation Graduate Research Fellowship Program under Grant No.

⁸For the models trained on FEWS, it is possible they have seen closely related senses to those in the zero-shot subsets of from the Wiktionary sense inventory; however, these are represented differently in each dataset and correspond to different definitions to be encoded by the biencoder.

DGE-1762114. The authors thank Julian Michael, Maarten Sap, and the UW NLP group for helpful conversations and comments on the work.

References

- Pierpaolo Basile, Annalina Caputo, and Giovanni Semeraro. 2014. An enhanced lesk word sense disambiguation algorithm through a distributional semantic model. In *Proceedings of COLING 2014, the 25th International Conference on Computational Linguistics: Technical Papers*, pages 1591–1600.
- Michele Bevilacqua and Roberto Navigli. 2020. Breaking through the 80% glass ceiling: Raising the state of the art in word sense disambiguation by incorporating knowledge graph information. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*.
- Terra Blevins and Luke Zettlemoyer. 2020. Moving down the long tail of word sense disambiguation with gloss informed bi-encoders. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186.
- Verena Henrich, Erhard Hinrichs, and Tatiana Vodolazova. 2012. Webcage: a web-harvested corpus annotated with germanet senses. In *Proceedings of the 13th Conference of the European Chapter of the Association for Computational Linguistics*, pages 387–396. Association for Computational Linguistics.
- Luyao Huang, Chi Sun, Xipeng Qiu, and Xuan-Jing Huang. 2019. Glossbert: Bert for word sense disambiguation with gloss knowledge. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3500–3505.
- Adam Kilgarriff. 2004. How dominant is the commonest sense of a word? In *International conference on text, speech and dialogue*, pages 103–111. Springer.
- Adam Kilgarriff and Joseph Rosenzweig. 2000. Framework and results for english senseval. *Computers and the Humanities*, 34(1-2):15–48.
- Diederik P. Kingma and Jimmy Lei Ba. 2015. [Adam: A method for stochastic optimization](#). In *Proceedings of 3rd International Conference of Learning Representations*.
- Sawan Kumar, Sharmistha Jat, Karan Saxena, and Partha Talukdar. 2019. [Zero-shot word sense disambiguation using sense definition embeddings](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 5670–5681, Florence, Italy. Association for Computational Linguistics.
- Daniel Loureiro and Jose Camacho-Collados. 2020. [Don’t neglect the obvious: On the role of unambiguous words in word sense disambiguation](#). *arXiv preprint arXiv:2004.14325*.
- Michael Matuschek and Iryna Gurevych. 2013. [Dijkstra-wsa: A graph-based approach to word sense alignment](#). *Transactions of the Association for Computational Linguistics*, 1:151–164.
- Christian M Meyer and Iryna Gurevych. 2011. What psycholinguists know about chemistry: Aligning wiktionary and wordnet for increased domain coverage. In *Proceedings of 5th International Joint Conference on Natural Language Processing*, pages 883–892.
- Rada Mihalcea. 2007. [Using Wikipedia for automatic word sense disambiguation](#). In *Human Language Technologies 2007: The Conference of the North American Chapter of the Association for Computational Linguistics; Proceedings of the Main Conference*, pages 196–203, Rochester, New York. Association for Computational Linguistics.
- George A Miller, Claudia Leacock, Randee Teng, and Ross T Bunker. 1993. A semantic concordance. In *Proceedings of the workshop on Human Language Technology*, pages 303–308. Association for Computational Linguistics.
- Tristan Miller and Iryna Gurevych. 2014. [WordNet—Wikipedia—Wiktionary: Construction of a three-way alignment](#). In *Proceedings of the Ninth International Conference on Language Resources and Evaluation (LREC’14)*, pages 2094–2100, Reykjavik, Iceland. European Language Resources Association (ELRA).
- Andrea Moro and Roberto Navigli. 2015. Semeval-2015 task 13: Multilingual all-words sense disambiguation and entity linking. In *Proceedings of the 9th international workshop on semantic evaluation (SemEval 2015)*, pages 288–297.
- Roberto Navigli. 2009. Word sense disambiguation: A survey. *ACM computing surveys (CSUR)*, 41(2):10.
- Roberto Navigli, David Jurgens, and Daniele Vannella. 2013. Semeval-2013 task 12: Multilingual word sense disambiguation. In *Second Joint Conference on Lexical and Computational Semantics (*SEM), Volume 2: Proceedings of the Seventh International Workshop on Semantic Evaluation (SemEval 2013)*, pages 222–231.
- Roberto Navigli and Simone Paolo Ponzetto. 2012. BabelNet: The automatic construction, evaluation and application of a wide-coverage multilingual semantic network. *Artificial Intelligence*, 193:217–250.

Martha Palmer, Christiane Fellbaum, Scott Cotton, Lauren Delfs, and Hoa Trang Dang. 2001. English tasks: All-words and verb lexical sample. In *Proceedings of SENSEVAL-2 Second International Workshop on Evaluating Word Sense Disambiguation Systems*, pages 21–24.

Jeffrey Pennington, Richard Socher, and Christopher Manning. 2014. [Glove: Global vectors for word representation](#). In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543. Association for Computational Linguistics.

Jason Phang, Thibault Févry, and Samuel R Bowman. 2018. [Sentence encoders on stilts: Supplementary training on intermediate labeled-data tasks](#). *arXiv preprint arXiv:1811.01088*.

Marten Postma, Ruben Izquierdo Bevia, and Piek Vossen. 2016. More is not always better: balancing sense distributions for all-words word sense disambiguation. In *Proceedings of COLING 2016, the 26th International Conference on Computational Linguistics: Technical Papers*, pages 3496–3506.

Sameer Pradhan, Edward Loper, Dmitriy Dligach, and Martha Palmer. 2007a. Semeval-2007 task-17: English lexical sample, srl and all words. In *Proceedings of the fourth international workshop on semantic evaluations (SemEval-2007)*, pages 87–92.

Sameer S Pradhan, Eduard Hovy, Mitch Marcus, Martha Palmer, Lance Ramshaw, and Ralph Weischedel. 2007b. Ontonotes: A unified relational semantic representation. In *International Conference on Semantic Computing (ICSC 2007)*, pages 517–526. IEEE.

Alessandro Raganato, Jose Camacho-Collados, and Roberto Navigli. 2017. Word sense disambiguation: A unified evaluation framework and empirical comparison. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 1, Long Papers*, pages 99–110.

Vincent Segonne, Marie Candito, and Benoît Crabbé. 2019. [Using Wiktionary as a resource for WSD : the case of French verbs](#). In *Proceedings of the 13th International Conference on Computational Semantics - Long Papers*, pages 259–270, Gothenburg, Sweden. Association for Computational Linguistics.

Benjamin Snyder and Martha Palmer. 2004. The english all-words task. In *Proceedings of SENSEVAL-3, the Third International Workshop on the Evaluation of Systems for the Semantic Analysis of Text*.

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, R’emi Louf, Morgan Funtowicz, and Jamie Brew. 2019. Huggingface’s transformers: State-of-the-art natural language processing. *ArXiv*, abs/1910.03771.

A Model Hyperparameters

Each model reported in this paper was tuned on a single hyperparameter sweep over the reported ranges and chosen based on the appropriate development set metric (accuracy on FEWS, F1 performance on the Unified WSD Framework).

Probe_{BERT} The linear layer in the BERT probe baseline is trained for 100 epochs. It is tuned over a range learning rates ($[5e-6, 1e-5, 5e-5, 1e-4]$, with a final learning rate of $1e-4$). We use a batch size of 128 to train this probe.

BEM For the biencoder model (BEM), we use the codebase provided by (Blevins and Zettlemoyer, 2020). Following this work, we train the BEM for 20 epochs with a warmup phase of 10,000 steps; we use a context batch size of 4 and a gloss batch size of 256. Each BEM is tuned over the following learning rates: $[1e-6, 5e-6, 1e-5, 5e-5]$. The BEM_{BERT} and BEM_{SemCor} had a final learning rate of $5e-6$, and the BEM_{FEWS}, of $1e-6$.

B Extended Train Baselines

The extended train set in FEWS contains all of the quotation-based examples from the train set as well as the additional, shorter example *illustrations* that are written by Wiktionary editors to exemplify a particular sense of a word. We retrain the **MFS**, **Probe_{BERT}**, and **BEM_{BERT}** baselines on this extended training set; the other baselines we consider (**Lesk** and **Lesk+emb**) are calculated without using either of the training sets.

Table 7 compare the extended train baselines against those trained on the standard train set. For the **MFS** and **Probe_{BERT}** baselines, we find that adding the extra, stylistically different *illustrations* in extended train slightly hurts performance. However, the stronger **BEM_{BERT}** is able to better use this data and achieves somewhat stronger performance with additional training data. Notably, most of this improvement in the **BEM_{BERT}** comes from the zero-shot evaluation setting, even though the extended train set does not contain any of these zero-shot senses.

		Dev			Test		
		Full Set	Few-shot	Zero-shot	Full Set	Few-shot	Zero-shot
MFS	Ext. Train	26.13	52.26	0.00	25.65	51.30	0.00
	Train	26.39	52.78	0.00	25.73	51.46	0.00
	Δ	-0.26	-0.52	–	-0.08	-0.16	–
Probe_{BERT}	Ext. Train	36.03	72.06	0.00	35.86	71.72	0.00
	Train	36.17	72.34	0.00	36.07	72.14	0.00
	Δ	-0.14	-0.28	–	-0.21	-0.42	–
BEM_{BERT}	Ext. Train	74.12	79.38	68.86	73.82	79.70	67.94
	Train	73.81	79.28	68.34	72.77	79.06	66.48
	Δ	0.31	0.10	0.52	1.05	0.64	1.46

Table 7: Accuracy of the FEWS baselines trained on the extended train set. In each group of rows, we report (1) the extended train baseline, (2) the comparable baseline trained on the standard train set, and (3) the performance delta between the two (where a positive delta indicates the extended train baseline performs better).